

A Flow-based Generative Network for Speech Synthesis

 BSD-3-Clause license

 2.3k stars

 530 forks

 Branches

 Tags

 Activity

 Star

 Notifications

 Code

 Issues 72

 Pull requests 7

 Actions

 Projects

 Security

 Insights

 ...

 5 Branches

 0 Tags





 Go to file

Go to file

Code



 **rafaelvalle** convert\_model.py: for compatibility with new pytorch 8afb643 · 4 years ago 

 tacotron2 @ 185cd24	tacotron2: updating tacotron2	4 years ago
 .gitmodules	waveglow	6 years ago
 LICENSE	LICENSE: adding license file	6 years ago
 README.md	README.md: reporting correct nu...	4 years ago
 config.json	config.json: adding optional tensor...	5 years ago
 convert_model.py	convert_model.py: for compatibility...	4 years ago
 denoiser.py	denoiser.py: adding code to remov...	5 years ago
 distributed.py	waveglow	6 years ago
 glow.py	Fused cond layers into single layer	5 years ago
 glow_old.py	Resolve PyTorch deprecation warni...	6 years ago
 inference.py	inference.py: using amp for mixed ...	5 years ago
 mel2samp.py	mel2samp.py: adding fmin and fma...	6 years ago
 requirements.txt	torch in requirements	6 years ago
 train.py	train.py: offsetting tensorboard itera...	5 years ago
 waveglow_logo.png	waveglow	6 years ago



# WaveGlow: a Flow-based Generative Network for Speech Synthesis

Ryan Prenger, Rafael Valle, and Bryan Catanzaro

In our recent [paper](#), we propose WaveGlow: a flow-based network capable of generating high quality speech from mel-spectrograms. WaveGlow combines insights from [Glow](#) and [WaveNet](#) in order to provide fast, efficient and high-quality audio synthesis, without the need for auto-regression. WaveGlow is implemented using only a single network, trained using only a single cost function: maximizing the likelihood of the training data, which makes the training procedure simple and stable.

Our [PyTorch](#) implementation produces audio samples at a rate of 1200 kHz on an NVIDIA V100 GPU. Mean Opinion Scores show that it delivers audio quality as good as the best publicly available WaveNet implementation.

Visit our [website](#) for audio samples.

## Setup

1. Clone our repo and initialize submodule

```
git clone https://github.com/NVIDIA/waveglow.git
cd waveglow
git submodule init
git submodule update
```



2. Install requirements `pip3 install -r requirements.txt`

3. Install [Apex](#)

## Generate audio with our pre-existing model

1. Download our [published model](#)
2. Download [mel-spectrograms](#)
3. Generate audio `python3 inference.py -f <(ls mel_spectrograms/*.pt) -w waveglow_256channels.pt -o . --is_fp16 -s 0.6`

N.b. use `convert_model.py` to convert your older models to the current model with fused residual and skip connections.

# Train your own model

📖 README

📄 BSD-3-Clause license



## 2. Make a list of the file names to use for training/testing

```
ls data/*.wav | tail -n+10 > train_files.txt  
ls data/*.wav | head -n10 > test_files.txt
```



## 3. Train your WaveGlow networks

```
mkdir checkpoints  
python train.py -c config.json
```



For multi-GPU training replace `train.py` with `distributed.py`. Only tested with single node and NCCL.

For mixed precision training set `"fp16_run": true` on `config.json`.

## 4. Make test set mel-spectrograms

```
python mel2samp.py -f test_files.txt -o . -c config.json
```

## 5. Do inference with your network

```
ls *.pt > mel_files.txt
```



## Releases

No releases published

## Packages

No packages published

## Contributors 7



## Languages

● Python 100.0%